

## ▼ Import Libraries

```
#Import all libraries
#Install EMIST library, import datasets of letters, Matplotlib
!pip3 install emnist
from emnist import list_datasets
from emnist import extract_training_samples
import matplotlib as mpl
import os
import numpy as np #linear algebra
import pandas as pd
from pandas_profiling import ProfileReport
import matplotlib.pyplot as plt
%matplotlib inline
import torch
import torchvision
import tensorflow as tf
import torch
import torch.nn as nn
print(tf.__version__)

import torch
import cv2
import torchvision.transforms as transforms
from tensorflow.keras.layers import *
from tensorflow.keras.models import *
import tensorflow_datasets as tfds
# import numpy as np
# import torch.nn.functional as F
# from torchvision.datasets import EMNIST
# from torch.utils.data import DataLoader
# import torchvision.transforms as tt
# from torch.utils.data import random_split
# from torchvision.utils import make_grid
# from tensorflow.keras.datasets import mnist
# from tensorflow.keras.utils import to_categorical
# from tensorflow.keras.models import Sequential
# from tensorflow.keras.layers import Conv2D
# from tensorflow.keras.layers import MaxPooling2D
# from tensorflow.keras.layers import Dense
# from tensorflow.keras.layers import Flatten
# from tensorflow.keras.optimizers import SGD

from PIL import Image
!pip install pyyaml h5py
import os

# from tensorflow import keras
# !pip install extra_keras_datasets
# from extra_keras_datasets import emnist

!pip install wandb
import wandb
from wandb.keras import WandbCallback
wandb.login()
```

```
!pip install tqdm
```

```
from numpy import argmax  
from keras.preprocessing.image import load_img  
from keras.preprocessing.image import img_to_array  
from keras.models import load_model
```

```
!pip install tqdm  
from tqdm import tqdm
```

```
import matplotlib  
import matplotlib.pyplot as plt  
%matplotlib inline
```

```
matplotlib.rcParams['figure.facecolor'] = '#ffffff'
```

```

one-any.whl (7.3 kB)
tqdm in /usr/local/lib/python3.7/dist-packages (from emnist) (4.63.0)
requests in /usr/local/lib/python3.7/dist-packages (from emnist) (2.23.0)
numpy in /usr/local/lib/python3.7/dist-packages (from emnist) (1.21.5)
idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests->emnist) (2.10)
certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests->emnist) (2021.10)
urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests->emnist) (1.25.1)
chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests->emnist) (3.0.4)
emnist
0.0

```

```

pyyaml in /usr/local/lib/python3.7/dist-packages (3.13)
h5py in /usr/local/lib/python3.7/dist-packages (3.1.0)
numpy>=1.14.5 in /usr/local/lib/python3.7/dist-packages (from h5py) (1.21.5)
cached-property in /usr/local/lib/python3.7/dist-packages (from h5py) (1.5.2)

```

```
2.py3-none-any.whl (1.7 MB)
```

```
██████████ | 1.7 MB 3.9 MB/s
```

```
.0
```

```
.4.0-py2.py3-none-any.whl (9.0 kB)
```

```
nvthon-dateutil>=2.6.1 in /usr/local/lib/nvthon3.7/dist-packages (from wandb) (2.8.2)
```

```
#Download dataset and it is 536 MB
```

```
list_datasets()
```

```
Downloading emnist.zip: 43.9MB [00:01, 40.5MB/s]
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
```

```
<ipython-input-2-aa4b3a38056e> in <module>()
```

```
1 #Download dataset and it is 536 MB
```

```
----> 2 list_datasets()
```

```
⏏ 4 frames
```

```
/usr/local/lib/python3.7/dist-packages/urllib3/response.py in stream(self, amt,
decode_content)
```

```
493             yield line
```

```
494         else:
```

```
--> 495             while not is_fp_closed(self._fp):
```

```
496                 data = self.read(amt=amt, decode_content=decode_content)
```

```
497
```

```
KeyboardInterrupt:
```

```
SEARCH STACK OVERFLOW
```

```
██████████ | 63 kB 1.3 MB/s
```

```
images, labels = extract_training_samples('byclass')
```

```
images.shape
```

```
urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/nvthon3.7/dist-packages (from requests->emnist) (1.25.1)
```

## GPU Running

```
packages: pathtools
```

```
# Get the GPU device name.
```

```
def get_default_device():
```

```
    """Pick GPU if available, else CPU"""
```

```
    if torch.cuda.is_available():
```

```
        return torch.device('cuda')
```

```
    else:
```

```
        return torch.device('cpu')
```

```

def to_device(data, device):
    """Move tensor(s) to chosen device"""
    if isinstance(data, (list,tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)

class DeviceDataLoader():
    """Wrap a dataloader to move data to a device"""
    def __init__(self, dl, device):
        self.dl = dl
        self.device = device

    def __iter__(self):
        """Yield a batch of data after moving it to device"""
        for b in self.dl:
            yield to_device(b, self.device)

    def __len__(self):
        """Number of batches"""
        return len(self.dl)

# Returns current working runtime

device = get_default_device()
device

```

## ▼ DataLoader Preparation

```

dataset = EMNIST(root="data/", split="byclass", download=True, train=True,
                 transform=tt.Compose([
                     lambda img: tt.functional.rotate(img, -90),
                     lambda img: tt.functional.hflip(img),
                     tt.ToTensor()
                 ]))

random_seed = 50
torch.manual_seed(random_seed);

val_size = 50000
train_size = len(dataset) - val_size

train_ds, val_ds = random_split(dataset, [train_size, val_size])
len(train_ds), len(val_ds)

batch_size = 400

train_dl = DataLoader(train_ds, batch_size, shuffle=True, num_workers=4, pin_memory=True)
val_dl = DataLoader(val_ds, batch_size*2, num_workers=4, pin_memory=True)

# Lets see a batch of images

def show_batch(dl):
    for images, labels in dl:
        fig, ax = plt.subplots(figsize=(12, 12))

```

```

fig, ax = plt.subplots(figsize=(12, 12))
ax.set_xticks([]); ax.set_yticks([])
ax.imshow(make_grid(images, nrow=20).permute(1, 2, 0))
break

```

```
show_batch(train_dl)
```

```

train_dl = DeviceDataLoader(train_dl, device)
val_dl = DeviceDataLoader(val_dl, device)

```

```
show_batch(train_dl_1)
```

# MNIST

## [MNIST recognition reference](#)

```

# save the final model to file
# load train and test dataset
def load_dataset(data):
    # load dataset
    (trainX, trainY), (testX, testY) = data
    # reshape dataset to have a single channel
    trainX = trainX.reshape((trainX.shape[0], 28, 28, 1))
    testX = testX.reshape((testX.shape[0], 28, 28, 1))
    # one hot encode target values
    trainY = to_categorical(trainY)
    testY = to_categorical(testY)
    return trainX, trainY, testX, testY

# scale pixels
def prep_pixels(train, test):
    # convert from integers to floats
    train_norm = train.astype('float32')
    test_norm = test.astype('float32')
    # normalize to range 0-1
    train_norm = train_norm / 255.0
    test_norm = test_norm / 255.0
    # return normalized images
    return train_norm, test_norm

# define cnn model
def define_model():
    model = Sequential()
    # (2,2) size of pooling area for max pooling    #(3,3) convolution kernel size    #32 number of convolutions
    model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform', input_shape=(28, 28, 1)))
    model.add(MaxPooling2D((2, 2)))
    model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform'))
    model.add(Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform'))
    model.add(MaxPooling2D((2, 2)))
    model.add(Flatten())
    model.add(Dense(100, activation='relu', kernel_initializer='he_uniform'))
    model.add(Dense(62, activation='softmax'))

```

```

# compile model
opt = SGD(learning_rate=0.01, momentum=0.9)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
return model

# run the test harness for evaluating a model
def run_test_harness(data):
    # load dataset

    trainX, trainY, testX, testY = load_dataset(data)
    # prepare pixel data
    trainX, testX = prep_pixels(trainX, testX)
    # define model
    model = define_model()
    # fit model
    model.fit(trainX, trainY, epochs=10, batch_size=32, verbose=0)
    # save model
    model.save('/content/sample_data/final_model.h5')

# entry point, run the test harness
data=mnist.load_data()
run_test_harness(data)

```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-2-61db1c596acd> in <module>()
      1 # entry point, run the test harness
----> 2 data=mnist.load_data()
      3 run_test_harness(data)

NameError: name 'mnist' is not defined

```

SEARCH STACK OVERFLOW

## ▼ SAVE MODEL

```

# load and prepare the image
def load_image(filename):
    # load the image
    img = load_img(filename, color_mode='grayscale', target_size=(28, 28))
    # convert to array
    img = img_to_array(img)
    # reshape into a single sample with 1 channel
    img = img.reshape(1, 28, 28, 1)
    # prepare pixel data
    img = img.astype('float32')
    img = img / 255.0
    return img

# load an image and predict the class
def run_example():
    # load the image
    img = load_image('/content/sample_data/img.png')

```

```

# load model
model = load_model('/content/sample_data/final_model.h5')
# predict the class
predict_value = model.predict(img)
digit = argmax(predict_value)
print(digit)

# entry point, run the example

print('RESULT:')
run_example()

```

## EMNIST DATASET RECOGNITION

[Reference](#)

[EMNIST Classification](#)

[EMNIST Classification NB](#)

[research Paper](#)

[CNN Theory](#)

[SAVE AND RESTORE MODEL FROM WANDB](#)

[github alphabet recognition](#)

TO FIND OUT WHAT IS TESTLOADER AND HOW TO GIVE FRAME BY FRAME IN VALIDATION DATA? or  
 emnist classification predicts on one image

RESEARCH FOR AROUBA AND SAIMA

### Points to note:

- We are using bymerge variant of EMNIST dataset. Here labels like j, o, i etc which look like J, O, I are merged.
- The EMNIST images provided here are inverted horizontally and rotated 90 anti-clockwise. For the ease of experimentation, we don't want to use it in this configuration. Thus we will rotate the image back by 90 deg anti-clockwise.
- We have a total of 814255 images. They are 28x28 pixels in resolution with only one channel.
- We have 697932 images as training data.
- We have 116323 images as testing data.
- We have 47 classes as shown:

```

# Gather EMNIST/bymerge dataset
train_ds, validation_ds = tfds.load(
    "emnist/bymerge",
    split=["train[:85%]", "train[85%:]"],
    as_supervised=True
)

```

#47 classes

```
LABELS = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9',  
          'A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q', 'R',  
          'a', 'b', 'd', 'e', 'f', 'g', 'h', 'n', 'q', 'r', 't']  
len(LABELS)
```

**Downloading and preparing dataset emnist/bymerge/3.0.0 (download: 535.73 MiB, generated: Unkno**

DI Completed...: 100% 1/1 [00:17<00:00, 12.60s/ url]

DI Size...: 100% 535/535 [00:17<00:00, 51.09 MiB/s]

Extraction completed...: 100% 1/1 [00:17<00:00, 17.55s/ file]

Extraction completed...: 100% 4/4 [00:12<00:00, 3.57s/ file]

Shuffling and writing examples to /root/tensorflow\_datasets/emnist/bymerge/3.0.0.incompleteGZV  
100% 697931/697932 [00:03<00:00, 288506.06 examples/s]

Shuffling and writing examples to /root/tensorflow\_datasets/emnist/bymerge/3.0.0.incompleteGZV  
100% 116322/116323 [00:00<00:00, 230190.69 examples/s]

**Dataset emnist downloaded and prepared to /root/tensorflow\_datasets/emnist/bymerge/3.0.0. Subs**  
47

```
AUTO = tf.data.experimental.AUTOTUNE  
BATCH_SIZE = 256
```

## We are transposing to rotate the image by 90 deg clockwise making the images human friendly.

```
def transpose_and_flatten(image, label=None):  
    image = tf.image.convert_image_dtype(image, dtype=tf.float32) # scale image pixels to [0,1]  
    image = tf.transpose(image, [1,0,2]) # transpose to get human friendly image, since rotation  
    image = tf.reshape(image, shape=(784,)) # permutation invariant or flatten  
  
    label = tf.one_hot(label, depth=len(LABELS)) # one hot encode label  
  
    return image, label
```

```
trainloader = (  
    train_ds  
    .shuffle(1024)  
    .map(transpose_and_flatten, num_parallel_calls=AUTO)  
    .batch(BATCH_SIZE)  
    .prefetch(AUTO)  
)
```

```
testloader = (  
    validation_ds  
    .map(transpose_and_flatten, num_parallel_calls=AUTO)  
    .batch(BATCH_SIZE)  
    .prefetch(AUTO)  
)
```


```
testloader
```



<PrefetchDataset element\_spec=(TensorSpec(shape=(None, 784), dtype=tf.float32, name=None), Ten

```
imgs, labels = next(iter(trainloader))
```

```
plt.figure(figsize=(10, 10))
for n in range(25):
    ax = plt.subplot(5, 5, n+1)
    plt.imshow(tf.reshape(imgs[n], shape=(28,28)), cmap='gray')
    plt.title(LABELS[np.argmax(labels[n])])
    plt.axis('off')
```



```
type(imgs[0])
```

```
tensorflow.python.framework.ops.EagerTensor
```

```
def DenseModel():
    inputs = Input(shape=(784,))
    x = Dense(256, activation='relu')(inputs)
    x = Dense(128, activation='relu')(x)
    outputs = Dense(len(LABELS), activation='softmax')(x) #Index recieved from LABELS list

    return Model(inputs=inputs, outputs=outputs)
```

```
early_stopper = tf.keras.callbacks.EarlyStopping(
```

```
monitor='val_loss', patience=5, verbose=0, mode='auto', restore_best_weights=True
)
```

```
# initialize wandb run
# do not change entity
wandb.init(entity='iit-bhu', project='emnist')

# hyperparameters
config = wandb.config
config.epochs = 70
config.learning_rate = 0.001

# model
tf.keras.backend.clear_session()
model = DenseModel()

# optimizer
optimizer = tf.keras.optimizers.Adam(learning_rate=config.learning_rate)

# compile
model.compile(optimizer, 'categorical_crossentropy', metrics=['acc'])

# train
model.fit(trainloader,
          epochs=config.epochs,
          validation_data=testloader,
          callbacks=[WandbCallback(),
                    early_stopper])
```

**wandb:** Currently logged in as: **insi29** (use `wandb login --relogin` to force relogin)

Tracking run with wandb version 0.12.11

Run data is saved locally in /content/wandb/run-20220304\_152316-1n4moujo

Syncing run **fluent-wave-243** to [Weights & Biases](#) ([docs](#))

Epoch 1/70

2318/2318 [=====] - 157s 67ms/step - loss: 0.6043 - acc: 0.8129 - val

Epoch 2/70

2318/2318 [=====] - 160s 69ms/step - loss: 0.3935 - acc: 0.8655 - val

Epoch 3/70

2318/2318 [=====] - 157s 68ms/step - loss: 0.3550 - acc: 0.8757 - val

Epoch 4/70

2318/2318 [=====] - 156s 67ms/step - loss: 0.3326 - acc: 0.8820 - val

Epoch 5/70

2318/2318 [=====] - 156s 67ms/step - loss: 0.3167 - acc: 0.8864 - val

Epoch 6/70

2318/2318 [=====] - 166s 72ms/step - loss: 0.3043 - acc: 0.8898 - val

Epoch 7/70

2318/2318 [=====] - 160s 69ms/step - loss: 0.2947 - acc: 0.8927 - val

Epoch 8/70

2318/2318 [=====] - 174s 75ms/step - loss: 0.2865 - acc: 0.8948 - val

Epoch 9/70

2318/2318 [=====] - 168s 73ms/step - loss: 0.2789 - acc: 0.8969 - val

Epoch 10/70

2318/2318 [=====] - 182s 79ms/step - loss: 0.2730 - acc: 0.8984 - val

<keras.callbacks.History at 0x7fd2f0efba10>

## ▼ SAVE AND RESTORE FROM WANDB

```
api = wandb.Api()

run = api.run("iit-bhu/emnist/1n4moujo")

model.save(os.path.join(wandb.run.dir, "model.h5"))

# Save a model file manually from the current directory:
wandb.save('model.h5')
```

```
[ '/content/wandb/run-20220304_152316-1n4moujo/files/model.h5' ]
```

```
# restore the model file "model.h5" from a specific run by user "lavanyashukla"
# in project "save_and_restore" from run "10pr4joa"
best_model = wandb.restore('model.h5', run_path="iit-bhu/emnist/1n4moujo")

# use the "name" attribute of the returned object
# if your framework expects a filename, e.g. as in Keras
model.load_weights(best_model.name)
```

## ▼ SAVE AND LOAD MODEL

```
model.save('/content/sample_data/final_model.h5')
model = load_model('/content/sample_data/final_model.h5')
```

## ▼ SAVE AND LOAD MODEL FROM DRIVE

```
!pip install pyyaml h5py
from google.colab import drive
drive.mount('/content/gdrive')
```

```
Requirement already satisfied: pyyaml in /usr/local/lib/python3.7/dist-packages (3.13)
Requirement already satisfied: h5py in /usr/local/lib/python3.7/dist-packages (3.1.0)
Requirement already satisfied: cached-property in /usr/local/lib/python3.7/dist-packages (from h5py) (1.5.2)
Requirement already satisfied: numpy>=1.14.5 in /usr/local/lib/python3.7/dist-packages (from h5py) (1.19.5)
Mounted at /content/gdrive
```

```
model.save('/content/gdrive/MyDrive/FYP@10Pearls/final.h5')
```

```
model_best=load_model('/content/gdrive/MyDrive/FYP@10Pearls/final.h5')
```

## ▼ Prediciton of Test Data

```

y_test = []      #Labels already known/expected
y_preds = []     #Labels predicted # o/p index
data_p=[]
for imgs, labels in tqdm(testloader):
    y_test.extend(np.argmax(labels, axis=1))
    y_pred = model_best.predict(imgs)
    y_preds.extend(np.argmax(y_pred, axis=1))
for i in (y_preds):
    data_p.append(LABELS[i])

Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (4.63.0)
100%|██████████| 409/409 [01:03<00:00, 6.41it/s]

```

## ▼ Convert Image to Tensor

```

import cv2
test=[]
def infer_prec(img, img_size):
    img = tf.expand_dims(img, -1)      # from 28 x 28 to 28 x 28 x 1
    img = tf.divide(img, 255)          # normalize
    img = tf.image.resize(img,         # resize acc to the input
        [img_size, img_size])
    img = tf.reshape(img,              # reshape to add batch dimension
        [784])
    return img

img = cv2.imread('/content/sample_data/sample_image.png', 0)  # read image as gray scale
print(img.shape)        # (720, 1280)

img = infer_prec(img, 28) # call preprocess function
print(img.shape)
img=tf.image.convert_image_dtype(
    img, dtype=tf.float32, name=None
)
test.append(img)
test=np.array(test)

(1480, 1490)
(784,)

plt.imshow(tf.reshape(img, shape=(28,28)), cmap='gray')

```

<matplotlib.image.AxesImage at 0x7fd2ef57c190>



type(test)

numpy.ndarray



tested=tf.convert\_to\_tensor(test, dtype=tf.float32)



tested2=tf.convert\_to\_tensor(img\_1, dtype=tf.float32)



type(tested)

tensorflow.python.framework.ops.EagerTensor


y\_pred = model\_best.predict(tested)  
#y\_pred # probabilities

# get predicted label  
pred=tf.argmax(y\_pred, axis=-1).numpy() # array([8], dtype=int64)

pred  
  
array([7])

LABELS[pred[0]]  
  
'7'


df=pd.DataFrame({'Expected Label':y\_test,'Predicted Label':y\_preds,'LABELS':data\_p})  
df.head(10)

	Expected Label	Predicted Label	LABELS	
0	28	28	S	
1	6	6	6	
2	1	1	1	
3	25	33	X	
4	2	2	2	
5	1	1	1	
6	0	0	0	
7	9	9	9	
8	12	12	C	
9	5	5	5	

```
plt.figure(figsize=(10, 10))
ax = plt.subplot(2, 2, 1)
plt.imshow(tf.reshape(imgs2[0], shape=(28,28)), cmap='gray')
plt.title(LABELS[y_preds[0]])
plt.axis('off')
```

(-0.5, 27.5, 27.5, -0.5)

5



```
DF=pd.DataFrame({'Labels':LABELS})
```

DF

Labels



0	0
1	1
2	2
3	3
4	4
5	5
6	6
7	7
8	8
9	9
10	A
11	B
12	C
13	D
14	E
15	F
16	G
17	H
18	I
19	J
20	K
21	L
22	M
23	N
24	O
25	P
26	Q
27	R

LABELS[28]

'S'

30U

y\_preds[0]

28

33X

y\_test[0]

^^

## ▼ Updates in Code

»»

u

```

def cv2_imshow(img):
    plt.imshow(img, cmap='gray')
def predict_image(img):
    image = cv2.imread(img)
    gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    blurred = cv2.GaussianBlur(gray, (5, 5), 0)
    edged = cv2.Canny(blurred, 30, 150)
    cnts = cv2.findContours(edged.copy(), cv2.RETR_EXTERNAL,
        cv2.CHAIN_APPROX_SIMPLE)
    cnts = imutils.grab_contours(cnts)
    cnts = sort_contours(cnts, method="left-to-right")[0]
    chars = []
    #    print(cnts)

    #cv2_imshow(edged)

    for c in cnts:
        # compute the bounding box of the contour
        (x, y, w, h) = cv2.boundingRect(c)
        # filter out bounding boxes, ensuring they are neither too small
        # nor too large
        if (w >= 5 and w <= 150) and (h >= 15 and h <= 120):
            # extract the character and threshold it to make the character
            # appear as *white* (foreground) on a *black* background, then
            # grab the width and height of the thresholded image
            roi = gray[y:y + h, x:x + w]
            thresh = cv2.threshold(roi, 0, 255,
                cv2.THRESH_BINARY_INV | cv2.THRESH_OTSU)[1]
            (tH, tW) = thresh.shape
            # if the width is greater than the height, resize along the
            # width dimension
            if tW > tH:
                thresh = imutils.resize(thresh, width=32)
            # otherwise, resize along the height
            else:
                thresh = imutils.resize(thresh, height=32)
            (tH, tW) = thresh.shape
            dX = int(max(0, 32 - tW) / 2.0)
            dY = int(max(0, 32 - tH) / 2.0)
            # pad the image and force 32x32 dimensions
            padded = cv2.copyMakeBorder(thresh, top=dY, bottom=dY,
                left=dX, right=dX, borderType=cv2.BORDER_CONSTANT,
                value=(0, 0, 0))
            padded = cv2.resize(padded, (28, 28))
            padded = padded.astype("float32")
            padded = np.expand_dims(padded, axis=-1)    #originally -1
            chars.append((padded, (x, y, w, h)))
    #    cv2_imshow(padded)

    #print("Padded: -")
    #cv2_imshow(padded)

```



```

"""
padded = cv2.copyMakeBorder(blurred, top=dY, bottom=dY,
                             left=dX, right=dX, borderType=cv2.BORDER_CONSTANT,
                             value=(0, 0, 0))
padded = cv2.resize(padded, (32, 32))
"""

# final_image=np.expand_dims(padded,axis=2)
print("characters ",len(chars))
boxes = [b[1] for b in chars]
#cv2_imshow(chars)
chars = np.array([c[0] for c in chars], dtype="float32")
# OCR the characters using our handwriting recognition model
# for c in chars:
preds = model.predict(chars)
# define the list of label names
labelNames = "0123456789"
labelNames += "ABCDEFGHIJKLMNOPQRSTUVWXYZabcdefghijklmnopqrstuvwxyz"
labelNames = [l for l in labelNames]
output=""
fig=plt.figure()
for c in range(len(chars)):
    print(c)
    # ax=fig.add_subplot(1,len(chars),c+1)
    # ax.imshow(chars[c],cmap='gray') #####
#fig.show()

for (pred, (x, y, w, h)) in zip(preds, boxes):
    i = np.argmax(pred)
    # print("i is ",i)
    # print("label length ",len(labelNames))
    #print("prediction is ",len(pred))
    prob = pred[i]
    label = labelNames[i]
    output+=label

return output
predict_image('/content/drive/MyDrive/Kaggle/aro.jpg')

```

## NANO NETS

## ▼ TRAINED DATA WITH MODEL

### [NANO NET MODEL](#)

```

import requests
import json
url = 'https://app.nanonets.com/api/v2/OCR/Model/866d66dc-69d6-433c-82c8-6123cb2db3b6/LabelFile/'

data = {'file': open('/content/sample_data/opencv_frame_0.png', 'rb')}

response = requests.post(url, auth=requests.auth.HTTPBasicAuth('N1E3KM3daTBYb6TrHNU0JdHLMe1tSNLB',
res=json.loads(response.text)

res["result"][0]["prediction"][0]["ocr_text"]

```